

PROBLEM STATEMENT - TO PREDICT/IDENTIFY THE NAME OF THE BREED BASED ON THE IMAGE OF THE DOG.

OVERVIEW:

1. Image recognition and classification have successfully applied in various domains, such as face recognition and scene understanding of autonomous driving.
2. At present, human face identification is successfully used for authentication and security purposes in many applications.
3. Human face identification has been widely used for animal identification, too. In particular, for the case of companion animals.
4. Since there are more than 180 dog breeds, dog breed recognition can be an essential task in order to provide proper training and health treatment.
5. Previously, dog breed recognition is done by human experts. However, some dog breeds might be challenging to evaluate due to the lack of experts and the difficulty of breeds' patterns themselves. It also takes time to find a classifier capable of determining a dog's breed.
6. In this competition, the goal of this case study is to build a classifier capable of determining a dog's breed.
7. We have provided 1200 images of dogs from 120 different breeds, and a test set of images and their labels are provided in order to practice fine-grained image categorization.
8. Each image has a filename that is unique to it.
9. Each image has a filename that is unique to it.
10. The data set contains 120 breeds of dogs.
11. The data set is balanced, so each class has approximately the same number of images.
12. Multi-class log loss between the predicted probability and the observed target has been used as the metric.

```
In [1]: import requests  
import os  
import zipfile  
# Set the URL of the zip file  
url = "https://storage.googleapis.com/kaggle-datasets/dogbreedidentification/dog-breed-identification.zip"  
# Set the path where you want to save the zip file  
path = "/content/dog-breed-identification.zip"  
  
# Download the zip file  
response = requests.get(url)  
with open(path, "wb") as f:  
    f.write(response.content)
```

```
In [2]: # Extract the zip file  
extract = zipfile.ZipFile(path)  
extract.extractall()
```

```
In [3]: # Import required libraries  
import pandas as pd  
import numpy as np  
import tensorflow as tf  
from tensorflow import keras  
import pickle  
import os
```

```
In [4]: # Read the labels.csv file  
labels = pd.read_csv('labels.csv')  
labels.head()
```

```
In [5]: # Read the CSV file containing the list of files and their corresponding breed  
files = pd.read_csv('files.csv')
```

```
In [6]: # Create a DataFrame to store the file paths and labels  
df = pd.DataFrame(files, columns=['path', 'label'])
```

```
In [7]: # Sort the DataFrame by path  
df = df.sort_values('path')
```

```
In [8]: # Group the DataFrame by breed and count the number of images per breed  
grouped = df.groupby('label').size().reset_index(name='Count')
```

```
In [9]: # Sort the grouped DataFrame by Count in descending order  
grouped = grouped.sort_values('Count', ascending=False)
```

```
In [10]: # Plot the bar chart showing the count of each Dog Breed using bar plot in descending order  
plt.figure(figsize=(10,10))  
plt.bar(grouped['label'], grouped['Count'], color='blue')  
plt.xlabel('Dog Breed')  
plt.ylabel('Count')  
plt.title('Value Counts for each Dog Breed')
```

Value Counts for each Dog Breed



OBSEVATIONS:

1. Counts of each breed looks good and the dataset is nearly balanced.
2. The Scottish Deerhound Dog Breed is maximum in number.
3. The Eskimo Dog and Briard Dog Breeds are minimum in number.

```
In [11]: # Plotting few images from the train data provided.  
# Set the figure size to (10,10)  
for i in range(20):  
    plt.subplot(5,4,i+1)  
    image = np.array(files['path'][i])  
    img = Image.open(image)  
    plt.imshow(img)  
    plt.title(files['label'][i])  
    plt.show()
```



TRAIN TEST SPLIT:

```
In [12]: from sklearn.model_selection import train_test_split  
Train, validation = train_test_split(df, test_size=0.33, stratify = files['label'], random_state=42)
```

ENCODING TARGET LABELS:

```
In [13]: from tensorflow import preprocessing  
le = preprocessing.LabelEncoder()  
le.fit(df['label'].to_list())  
Train_image_labels = le.transform(df['label'].to_list())
```

```
In [14]: validation_image_labels = validation['label'].to_list()
```

```
In [15]: with open('/content/drive/MyDrive/dog_breed/encoder','wb') as file:  
    pickle.dump(le,file)
```

```
In [16]: with open('/content/drive/MyDrive/dog_breed/labels','wb') as file:  
    pickle.dump(list(le.classes_),file)
```

```
In [17]: Train_image_paths = Train['path'].tolist()  
validation_image_paths = validation['path'].tolist()
```

```
In [18]: def fdata_generator(images, labels, is_training, batch_size=16):  
    def parse_function(filename, labels):  
        image = tf.io.read_file(filename)  
        image = tf.image.decode_jpeg(image, channels=3)  
        image = tf.image.resize(image, [256, 256])  
        image = tf.image.convert_image_dtype(image, tf.float32)  
        image = tf.image.random_flip_left_right(image)  
        image = tf.image.random_flip_up_down(image)  
        image = tf.image.random_hue(image, 0.5)  
        image = tf.image.random_saturation(image, 0.5, 2.5)  
        image = tf.image.random_brightness(image, 0.5)  
        image = tf.image.random_contrast(image, 0.1, 2)  
        return image, y  
    def augment(image, labels):  
        image = tf.image.random_flip_left_right(image)  
        image = tf.image.random_flip_up_down(image)  
        image = tf.image.random_hue(image, 0.5)  
        image = tf.image.random_saturation(image, 0.5, 2.5)  
        image = tf.image.random_brightness(image, 0.5)  
        image = tf.image.random_contrast(image, 0.1, 2)  
        return image, labels  
    dataset = tf.data.Dataset.from_tensor_slices((images,labels))  
    if is_training:  
        dataset = dataset.map(augment)  
        dataset = dataset.batch(batch_size, drop_remainder=True)  
    else:  
        dataset = dataset.repeat()  
    dataset = dataset.prefetch(tf.data.experimental.AUTOTUNE)  
    return dataset
```

```
In [19]: train_generator = fdata_generator(Train_image_paths, Train_image_labels, is_training=True)
```

```
In [20]: validation_generator = fdata_generator(validation_image_paths, validation_image_labels, is_training=False)
```

```
In [21]: steps_per_epoch = np.ceil(len(Train_image_paths)/16)
```

```
In [22]: validation_steps = np.ceil(len(validation_image_paths)/16)
```

```
In [23]: from tensorflow import keras
```

```
In [24]: base_inception_v3 = InceptionV3(weights = 'imagenet', include_top = False, input_shape = input_shape)
```

```
In [25]: model_inception = GlobalAveragePooling2D()(base_inception_v3.output)
```

```
In [26]: model_resnet = ResNet50(weights = 'imagenet', include_top = False, input_shape = input_shape)
```

```
In [27]: model_inception = Model(inputs=base_inception_v3.input, outputs=model_inception)
```

```
In [28]: model_resnet = Model(inputs=model_resnet.input, outputs=model_resnet.output)
```

```
In [29]: merged = Concatenate()([model_inception.output, model_resnet.output])
```

```
In [30]: x = BatchNormalization()(merged)
```

```
In [31]: x = Dropout(0.5)(x)
```

```
In [32]: x = Dense(120, activation = 'softmax')(x)
```

```
In [33]: model = Model(inputs=base_inception_v3.input, outputs=x)
```

```
In [34]: model.compile(optimizer=keras.optimizers.SGD(lr=0.001, momentum=0.9), loss='categorical_crossentropy', metrics=['categorical_accuracy'])
```

```
In [35]: Model: "model"
```

```
In [36]: Layer (type) Output Shape Param # Connected to  
=====
```

```
In [37]: input_(InputLayer) [(None, 400, 400, 3) []]  
=====
```

```
In [38]: inception_v3 (Functional) [(None, 11, 11, 2048) 21802784 ['input_1[0][0]']]  
=====
```

```
In [39]: resnet152 (Functional) [(None, 13, 13, 2048) 58370944 ['input_1[0][0]']]  
=====
```

```
In [40]: global_average_pooling2d_1 (GlobalAveragePooling2D) [(None, 2048) 0 ['inception_v3[0][0]']]  
=====
```

```
In [41]: flatten (Flatten) [(None, 2048) 0 ['global_average_pooling2d_1[0][0]']]  
=====
```

```
In [42]: flatten_1 (Flatten) [(None, 2048) 0 ['global_average_pooling2d_1[0][0]']]  
=====
```

```
In [43]: concatenate_2 (Concatenate) [(None, 4096) 16384 ['concatenate_2[0][0]']]  
=====
```

```
In [44]: batch_normalization_94 (BatchNormalizat... 16384 ['batch_normalization_94[0][0]'])  
=====
```

```
In [45]: dense (Dense) [(None, 256) 1848832 ['batch_normalization_94[0][0]']]  
=====
```

```
In [46]: dropout (Dropout) [(None, 256) 0 ['dense[0][0]']]  
=====
```

```
In [47]: dense_3 (Dense) [(None, 120) 30840 ['dropout[0][0]']]  
=====
```

```
In [48]: Total params: 81,269,784  
Trainable params: 1,083,056  
Non-trainable params: 80,181,920
```

```
/usr/local/lib/python3.8/dist-packages/keras/optimizers/optimizer_v2/gradient_descent.py:108: UserWarning: The `lr` argument is deprecated; please use `learning_rate` instead.  
super(SGD, self).__init__(name, **kwargs)
```

```
In [49]: from keras.callbacks import ModelCheckpoint, ReduceLROnPlateau
```

```
S, reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.9, patience=2, verbose=0)
```

```
In [50]: Inception_resnet_stacked.fit_generator(train_generator, steps_per_epoch=steps_per_epoch, validation_data=validation_generator, validation_steps=validation_steps, callbacks=[reduce_lr])
```

```
In [51]: Epoch 1/10  
428/428 [=====] ETA: 0s - loss: 0.62699, saving model to /content/drive/MyDrive/dog_breed/Inception_resnet_stacked_weights.h5
```

```
In [52]: 428/428 [=====] 319s 67ms/step - loss: 0.2319 - categorical_accuracy: 0.5976 - val_loss: 0.62699 - val_categorical_accuracy: 0.5976
```

```
In [53]: Epoch 2/10  
428/428 [=====] ETA: 0s - loss: 0.6038 - categorical_accuracy: 0.5831 - val_loss: 0.6038 - val_categorical_accuracy: 0.5831
```

```
In [54]: Epoch 3/10  
428/428 [=====] ETA: 0s - loss: 0.5805 - categorical_accuracy: 0.5934 - val_loss: 0.5805 - val_categorical_accuracy: 0.5934
```

```
In [55]: Epoch 4/10  
428/428 [=====] ETA: 0s - loss: 0.5572 - categorical_accuracy: 0.6031 - val_loss: 0.5572 - val_categorical_accuracy: 0.6031
```

```
In [56]: Epoch 5/10  
428/428 [=====] ETA: 0s - loss: 0.5347 - categorical_accuracy: 0.6136 - val_loss: 0.5347 - val_categorical_accuracy: 0.6136
```

```
In [57]: Epoch 6/10  
428/428 [=====] ETA: 0s - loss: 0.5129 - categorical_accuracy: 0.6232 - val_loss: 0.5129 - val_categorical_accuracy: 0.6232
```

```
In [58]: Epoch 7/10  
428/428 [=====] ETA: 0s - loss: 0.4903 - categorical_accuracy: 0.6325 - val_loss: 0.4903 - val_categorical_accuracy: 0.6325
```

```
In [59]: Epoch 8/10  
428/428 [=====] ETA: 0s - loss: 0.4676 - categorical_accuracy: 0.6421 - val_loss: 0.4676 - val_categorical_accuracy: 0.6421
```

```
In [60]: Epoch 9/10  
428/428 [=====] ETA: 0s - loss: 0.4448 - categorical_accuracy: 0.6517 - val_loss: 0.4448 - val_categorical_accuracy: 0.6517
```

```
In [61]: Epoch 10/10  
428/428 [=====] ETA: 0s - loss: 0.4221 - categorical_accuracy: 0.6614 - val_loss: 0.4221 - val_categorical_accuracy: 0.6614
```

```
In [62]: Inception_resnet_stacked.compile(optimizer=keras.optimizers.SGD(lr=0.001, momentum=0.9), loss='categorical_crossentropy', metrics=['categorical_accuracy'])
```

```
In [63]: Inception_resnet_stacked.summary()
```

```
In [64]: Model: "Inception_resnet_stacked"
```

```
In [65]: Layer (type) Output Shape Param # Connected to  
=====
```

```
In [66]: input_(InputLayer) [(None, 400, 400, 3) []]  
=====
```

```
In [67]: inception_v3 (Functional) [(None, 11, 11, 2048) 21802784 ['input_1[0][0]']]  
=====
```

```
In [68]: resnet152 (Functional) [(None, 13, 13, 2048) 58370944 ['input_1[0][0]']]  
=====
```

```
In [69]: global_average_pooling2d_1 (GlobalAveragePooling2D) [(None, 2048) 0 ['inception_v3[0][0]']]  
=====
```

```
In [70]: flatten (Flatten) [(None, 2048) 0 ['global_average_pooling2d_1[0][0]']]  
=====
```

```
In [71]: flatten_1 (Flatten) [(None, 2048) 0 ['global_average_pooling2d_1[0][0]']]  
=====
```

```
In [72]: concatenate_2 (Concatenate) [(None, 4096) 16384 ['concatenate_2[0][0]']]  
=====
```

```
In [73]: batch_normalization_94 (BatchNormalizat... 16384 ['batch_normalization_94[0][0]'])  
=====
```

```
In [74]: dense (Dense) [(None, 256) 1848832 ['batch_normalization_94[0][0]']]  
=====
```

```
In [75]: dropout (Dropout) [(None, 256) 0 ['dense[0][0]']]  
=====
```

```
In [76]: dense_3 (Dense) [(None, 120) 30840 ['dropout[0][0]']]  
=====
```

```
In [77]: Total params: 81,269,784  
Trainable params: 1,083,056  
Non-trainable params: 80,181,920
```

```
/usr/local/lib/python3.8/dist-packages/keras/optimizers/optimizer_v2/gradient_descent.py:108: UserWarning: The `lr` argument is deprecated; please use `learning_rate` instead.  
super(SGD, self).__init__(name, **kwargs)
```

```
In [78]: from keras.callbacks import ModelCheckpoint, ReduceLROnPlateau
```

```
S, reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.9, patience=2, verbose=0)
```

```
In [79]: Inception_resnet_stacked.fit_generator(train_generator, steps_per_epoch=steps_per_epoch, validation_data=validation_generator, validation_steps=validation_steps, callbacks=[reduce_lr])
```

```
In [80]: Epoch 1/10  
428/428 [=====] ETA: 0s - loss: 0.62699, saving model to /content/drive/MyDrive/dog_breed/Inception_resnet_stacked_weights.h5
```

```
In [81]: 428/428 [=====] 319s 67ms/step - loss: 0.2319 - categorical_accuracy: 0.5976 - val_loss: 0.62699 - val_categorical_accuracy: 0.5976
```

```
In [82]: Epoch 2/10  
428/428 [=====] ETA: 0s - loss: 0.6038 - categorical_accuracy: 0.5831 - val_loss: 0.6038 - val_categorical_accuracy: 0.5831
```

```
In [83]: Epoch 3
```